

# Sheet#01 Proposed Solution MLPH\_W24

J. Schubert, M. Schümann

October 26, 2024

## 1 Sheet 1

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib import pyplot as plt
%matplotlib inline
```

### 1.1 1 Principal Component Analysis

#### 1.1.1 (a)

```
[2]: # TODO: implement PCA (fill in the blanks in the function below)

def pca(data, n_components=None):
    """
    Principal Component Analysis on a  $p \times N$  data matrix.

    Parameters
    -----
    data : np.ndarray
        Data matrix of shape  $(p, N)$ .
    n_components : int, optional
        Number of requested components. By default returns all components.

    Returns
    -----
    np.ndarray, np.ndarray
        the pca components (shape  $(n\_components, p)$ ) and the projection (shape  $\hookrightarrow (n\_components, N)$ )

    """
    # set n_components to p by default
    n_components = data.shape[0] if n_components is None else n_components
    assert n_components <= data.shape[0], f"Got n_components larger than  $\hookrightarrow$ 
    dimensionality of data!"

    # center the data
```

```

X_c = data - np.expand_dims(np.mean(data, axis=1), axis=-1)

# compute X times X transpose
XXT = X_c @ X_c.T
assert XXT.shape == (data.shape[0], data.shape[0])

# compute the eigenvectors and eigenvalues
from numpy.linalg import eig
evalues, evectors = eig(XXT)

# sort the eigenvectors by eigenvalue and take the n_components largest ones

evalues, evectors = zip(*sorted(zip(evalues, evectors), key=lambda x: x[0],
↪reverse=True))

# verify that the sorting went well / indeed all eigenvectors are shape (p,)
assert all(map(lambda x: x.shape == (data.shape[0],), evectors))

# compute X_projected, the projection of the data to the components
basis = np.array(evectors[:n_components])
X_projected = basis @ X_c

return basis, X_projected # return the n_components first components and
↪the pca projection of the data

# test
_, _ = pca(np.random.random(size=(100, 1000)));

```

```

[3]: # Example data to test your implementation
# All the asserts on the bottom should go through if your implementation is
↪correct

data = np.array([
    [ 1,  0,  0, -1,  0,  0],
    [ 0,  3,  0,  0, -3,  0],
    [ 0,  0,  5,  0,  0, -5]
], dtype=np.float32)

# add a random offset to all samples. it should not affect the results
data += np.random.randn(data.shape[0], 1)

n_components = 2
components, projection = pca(data, n_components=n_components) # apply your
↪implementation

# the correct results are known (up to some signs)

```

```

true_components = np.array([[0, 0, 1], [0, 1, 0]], dtype=np.float32)
true_projection = np.array([
    [ 0, 0, 5, 0, 0, -5],
    [ 0, 3, 0, 0, -3, 0]
], dtype=np.float32)

# check that components match, up to sign
assert isinstance(components, np.ndarray), f'Expected components to be numpy_
↳array but got {type(components)}'
assert components.shape == true_components.shape, f'{components.shape}!
↳={true_components.shape}'
assert np.allclose(np.abs(components * true_components).sum(1), np.
↳ones(n_components)), f'Components not matching'

# check that projections agree, taking into account potentially flipped_
↳components
assert isinstance(projection, np.ndarray), f'Expected projection to be numpy_
↳array but got {type(projection)}'
assert projection.shape == (n_components, data.shape[1]), f'Incorrect shape of_
↳projection: Expected {(n_components, data.shape[1])}, got {projection.shape}'
assert np.allclose(projection, true_projection * (components * true_components).
↳sum(1, keepdims=True), atol=1e-6), f'Projections not matching'

print('Test successful!')

```

Test successful!

### 1.1.2 (b)

Load the data (it is a subset of the data at <https://opendata.cern.ch/record/4910#>)

```

[4]: features = np.load('data/dijet_features.npy')
labels = np.load('data/dijet_labels.npy')
label_names = ['b', 'c', 'q'] # bottom, charm or light quarks

print(f'{features.shape=}, {labels.shape=}') # print the shapes

# TODO: print how many samples of each class are present in the data (hint:_
↳numpy.unique)
for label in np.unique(labels).astype(int):
    print(f"Label \"{label_names[label]}\": {(labels == label).sum()} total_
↳samples.")

```

```

features.shape=(116, 2233), labels.shape=(2233,)
Label "b": 999 total samples.
Label "c": 864 total samples.
Label "q": 370 total samples.

```

Normalize the data

```
[5]: # TODO: report range of features and normalize the data to zero mean and unit_
      ↪variance
      features = features / np.expand_dims(np.std(features, axis=1), axis=-1)
```

### 1.1.3 (c)

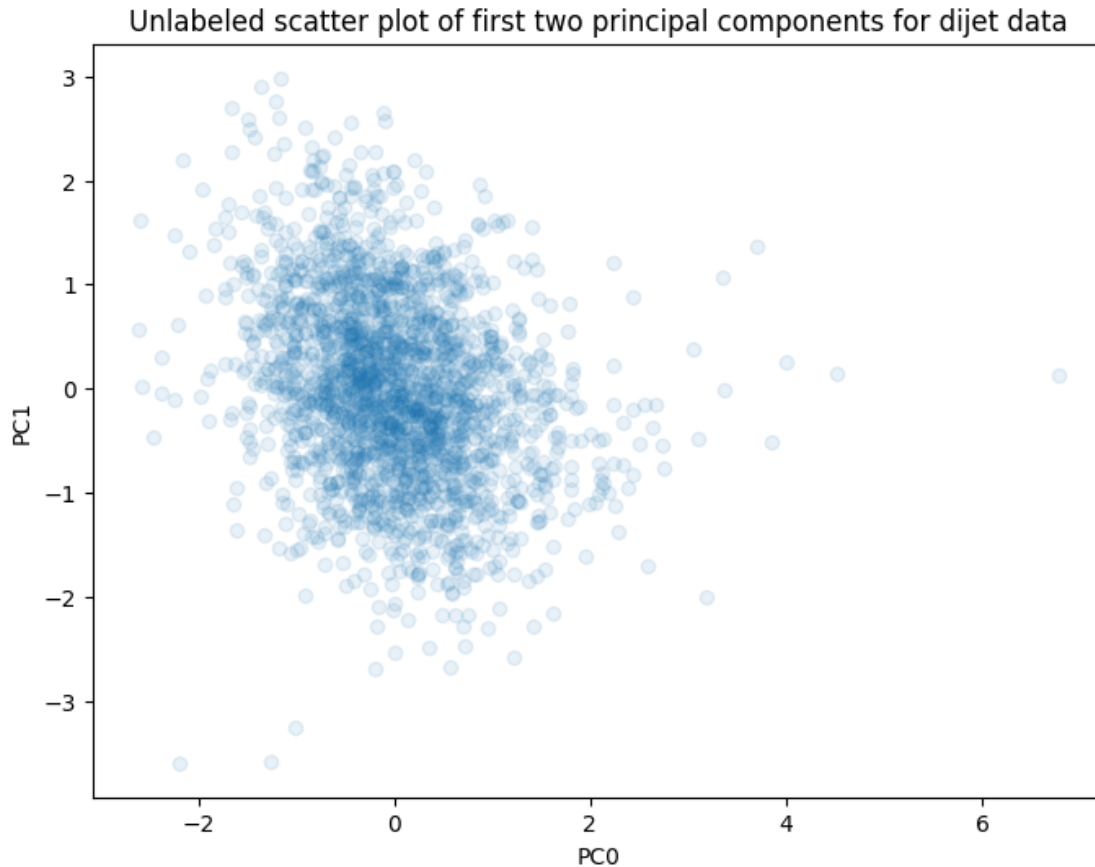
Compute a 2D PCA projection and make a scatterplot of the result, once without color, once coloring the dots by label. Interpret your results.

```
[6]: # TODO: apply PCA as implemented in (a)
      components, projection = pca(features, n_components=2)
```

```
[7]: # TODO: make a scatterplot of the PCA projection
      plt.rcParams["figure.figsize"] = 8,6
      plt.scatter(*projection, alpha=0.1)
      plt.xlabel("PC0")
      plt.ylabel("PC1")
      plt.title("Unlabeled scatter plot of first two principal components for dijet_
      ↪data")
```

```
/home/jona/.local/lib/python3.10/site-packages/matplotlib/cbook.py:1699:
ComplexWarning: Casting complex values to real discards the imaginary part
    return math.isfinite(val)
/home/jona/.local/lib/python3.10/site-packages/matplotlib/collections.py:194:
ComplexWarning: Casting complex values to real discards the imaginary part
    offsets = np.asanyarray(offsets, float)
```

```
[7]: Text(0.5, 1.0, 'Unlabeled scatter plot of first two principal components for
      dijet data')
```

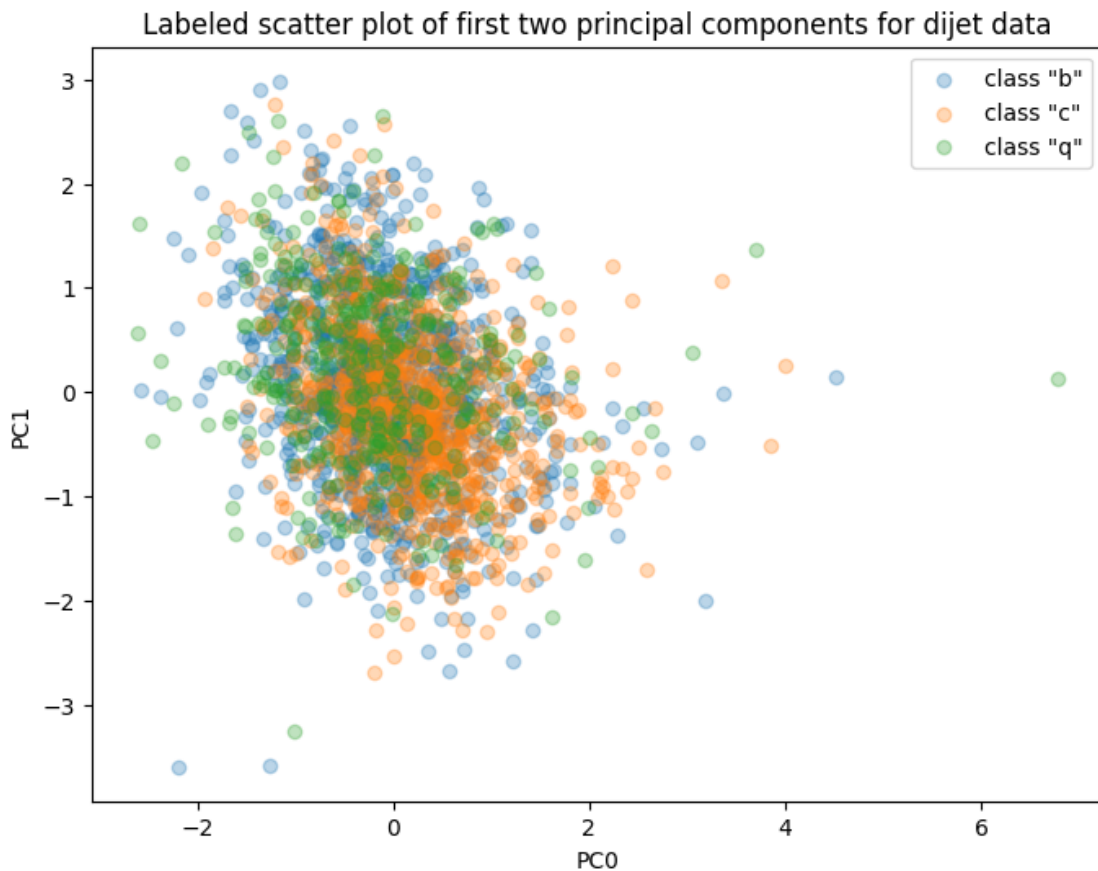


From this plot one can not tell whether the different classes are separated or not.

```
[8]: # TODO: make a scatterplot, coloring the dots by their label and including a
      ↪ legend with the label names
      # (hint: one way is to call plt.scatter once for each of the three possible
      ↪ labels. Why could it be problematic to scatter the data sorted by labels
      ↪ though?)
      for label in np.unique(labels).astype(int):
          mask = labels == label
          relevant_projection = projection[:, mask]
          plt.scatter(*relevant_projection, alpha=0.3, label=f"class
          ↪ \"{label_names[label]}\"")

      plt.legend()
      plt.xlabel("PC0")
      plt.ylabel("PC1")
      plt.title("Labeled scatter plot of first two principal components for dijet
      ↪ data")
```

```
[8]: Text(0.5, 1.0, 'Labeled scatter plot of first two principal components for dijet data')
```



In this plot one can see that the different classes are not well separated within the plane. This shows that PCA is a solid first step, but as seen in the lecture has rather lackluster results. In order to improve one has to use different (non-linear) methods.

## 1.2 2 Nonlinear Dimension Reduction

```
[9]: import umap # import umap-learn, see https://umap-learn.readthedocs.io/
```

```
/home/jona/.local/lib/python3.10/site-packages/tqdm/auto.py:21: TqdmWarning:  
IPProgress not found. Please update jupyter and ipywidgets. See  
https://ipywidgets.readthedocs.io/en/stable/user_install.html  
from .autonotebook import tqdm as notebook_tqdm
```

```
[10]: # if you have not done 1(b) yet, you can load the normalized features directly:  
features = np.load('data/dijet_features_normalized.npy')  
labels = np.load('data/dijet_labels.npy')  
label_names = ['b', 'c', 'q'] # bottom, charm or light quarks
```

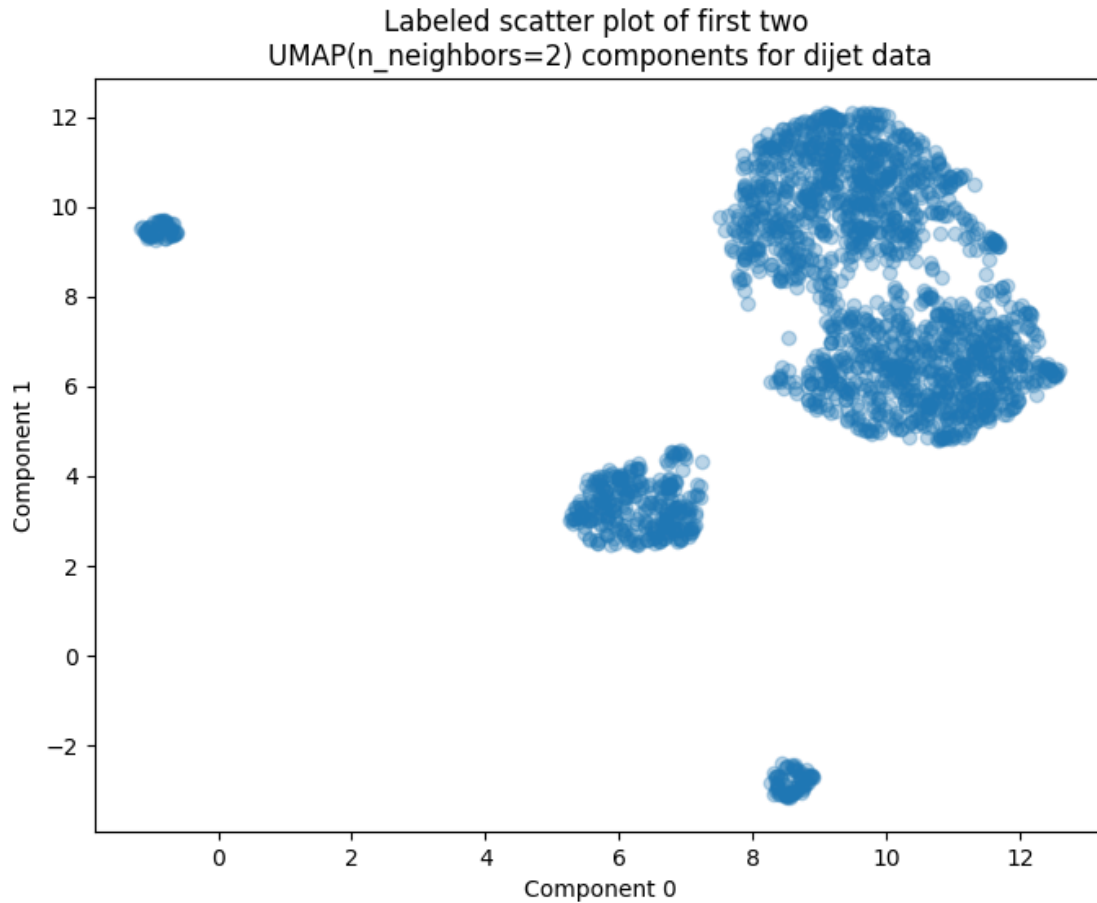
### 1.2.1 (a)

```
[11]: # TODO: Apply umap on the normalized jet features from exercise 1. It will
      ↪ take a couple of seconds.
      # note: umap uses a different convention regarding the feature- and sample
      ↪ dimension,  $N \times p$  instead of  $p \times N$ !

      n_neighbors = 2
      reducer = umap.UMAP(n_components=n_neighbors)

[15]: # TODO: make a scatterplot of the UMAP projection
      def create_scatter_plot(transformed_features, title=None):
          # Combine all data points into a single scatter plot
          plt.scatter(*transformed_features, alpha=0.3) # Set color for all points
          plt.xlabel("Component 0")
          plt.ylabel("Component 1")
          plt.title(title if title else "")


      # Assuming `reducer` and `features` are already defined and initialized
      projection = reducer.fit_transform(features.T).T
      create_scatter_plot(projection, title=f"Labeled scatter plot of first two
      ↪ \nUMAP(n_neighbors={n_neighbors}) components for dijet data")
```



Looking at this plot it seems that the three classes are well separated as we have four separate areas in the plane.

```
[12]: def create_scatter_plot(transformed_features, title=None):
    for label in np.unique(labels).astype(int):
        mask = labels == label
        relevant_projection = projection[:, mask]
        plt.scatter(*relevant_projection, alpha=0.3, label=f"class_{label}")
    plt.legend()
    plt.xlabel("component 0")
    plt.ylabel("component 1")
    plt.title(title if title else "")
    projection = reducer.fit_transform(features.T, ).T
    create_scatter_plot(projection, title=f"Labeled scatter plot of first two_
    ↪umap(n_neighbors={n_neighbors}) components for dijet data")
```





Looking at the colored in version we can see that b and c are not wwell seperated from each other but only have a small overlap with q.

### 1.2.2 (b)

```
[21]: for n_neighbors in (2, 4, 8, 15, 30, 60, 100):
    # TODO: repeat the above, varying the n_neighbors parameter of UMAP
    reducer = umap.UMAP(n_neighbors=n_neighbors)
    projection = reducer.fit_transform(features.T, ).T
    create_scatter_plot(projection, title=f"Labeled scatter plot of first two_
    ↪\numap(n_neighbors={n_neighbors}) components for dijet data")
    plt.show()
```

```
/root/miniconda3/envs/mlph/lib/python3.9/site-
packages/sklearn/manifold/_spectral_embedding.py:455: UserWarning: Exited at
iteration 2000 with accuracies
[2.93465312e-15 2.05340802e-06 3.17304847e-06 9.69502793e-06]
not reaching the requested tolerance 4.664063453674316e-06.
Use iteration 1955 instead with accuracy
```

2.2299514031715795e-06.

```
_, diffusion_map = lobpcg(  
/root/miniconda3/envs/mlph/lib/python3.9/site-  
packages/sklearn/manifold/_spectral_embedding.py:455: UserWarning: Exited  
postprocessing with accuracies  
[4.95009064e-15 1.32157027e-06 2.49654518e-06 5.10660318e-06]  
not reaching the requested tolerance 4.664063453674316e-06.  
_, diffusion_map = lobpcg(  
/root/miniconda3/envs/mlph/lib/python3.9/site-packages/umap/spectral.py:550:  
UserWarning: Spectral initialisation failed! The eigenvector solver  
failed. This is likely due to too small an eigengap. Consider  
adding some noise or jitter to your data.
```

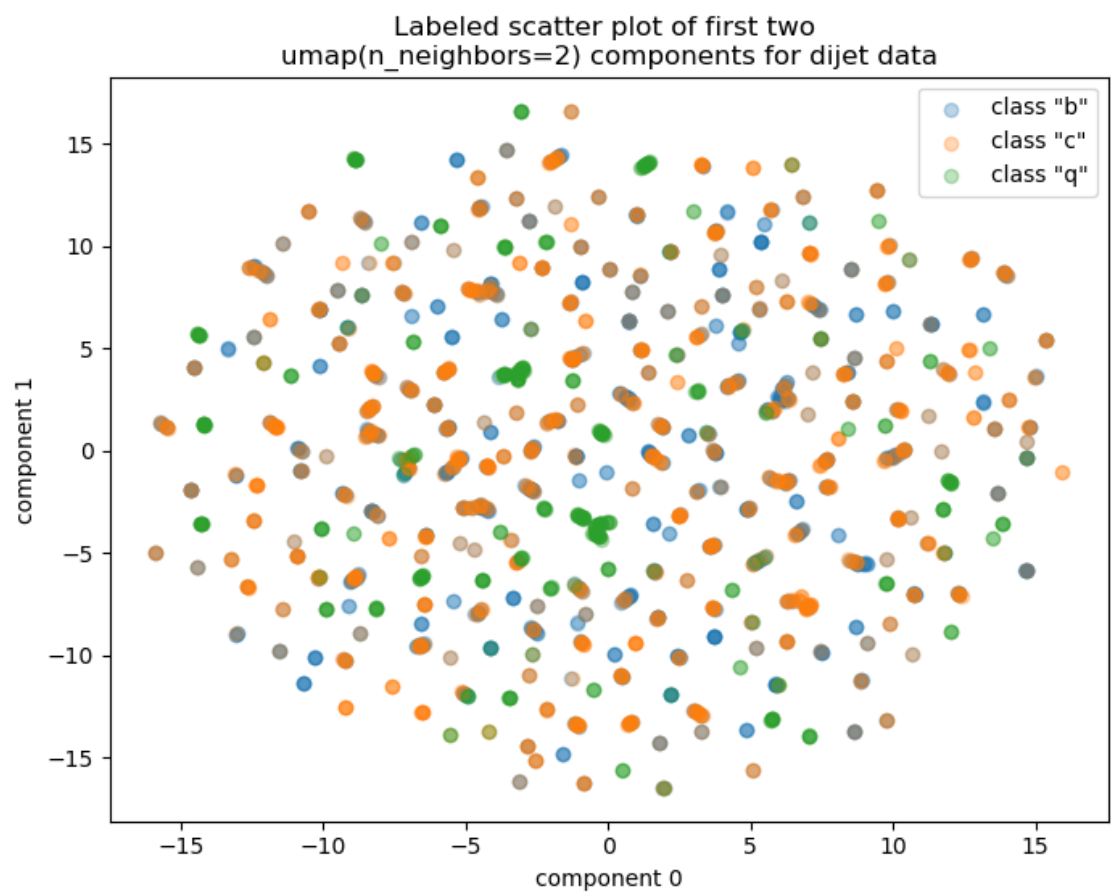
Falling back to random initialisation!

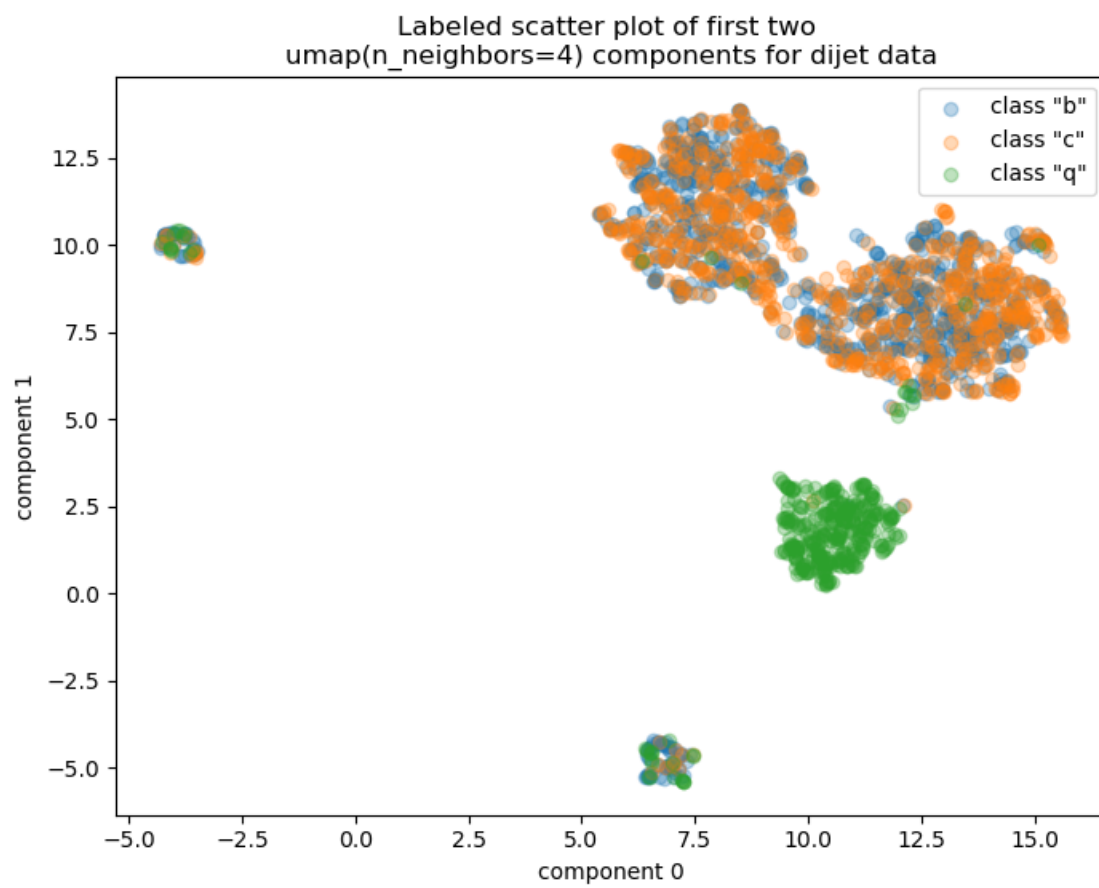
```
warn(  
/root/miniconda3/envs/mlph/lib/python3.9/site-packages/umap/spectral.py:550:  
UserWarning: Spectral initialisation failed! The eigenvector solver  
failed. This is likely due to too small an eigengap. Consider  
adding some noise or jitter to your data.
```


Falling back to random initialisation!


```
warn(  


```



















From these plots one can observe that the clustersizes increase with  $n$ . And the single points in a cluster are further appart. This allows us to chose  $n$  according to the task; if the goal is to visualize small, tight clusters, lower values are preferred, whereas understanding broader relationships benefits from higher values.

### 1.3 3 RANSAC

- $p$ : Probability of an inlier
- $1 - p$ : Probability of an outlier
- $m$ : Number of data points (small sample for outliers)
- With this we get:
  - $P(\text{inlier-free subset}) = p^m$
  - $P(\text{not inlier-free subset}) = (1 - p_m)$
  - $P(\text{not inlier-free subset in } r \text{ runs}) = (1 - p_m)^r$
  - $P(\text{at least one inlier-free subset}) = 1 - (1 - p_m)^r$

So, we get:

$$1 - (1 - p^m)^r \geq 0.99 \quad (1)$$

$$(1 - p^m)^r \leq 0.01 \quad (2)$$

$$\ln((1 - p^m)^r) \leq \ln(0.01) \quad (3)$$

Simplifying further:

$$r \geq \frac{\ln(0.01)}{\ln(1 - p^m)} = -\frac{2\ln(10)}{\ln(1 - p^m)} \quad (4)$$